

# TECHNOLOGY SERVICES AGREEMENT

Contract ID: TECH-2026-001

THIS AGREEMENT is entered into as of **January 15, 2026** (the "Effective Date") by and between:

| Service Provider                                                                    | Client                                                                   |
|-------------------------------------------------------------------------------------|--------------------------------------------------------------------------|
| NovaTech Solutions Ltd.<br>1 Bay Street, Suite 800<br>Toronto, ON M5J 2N8<br>Canada | Global Industries LLC<br>250 Bay Street<br>Toronto, ON M5J 2T3<br>Canada |

**WHEREAS**, NovaTech Solutions Ltd. provides cloud infrastructure and AI integration services;

**WHEREAS**, Global Industries LLC desires to engage NovaTech for an AI-powered document processing pipeline;

**WHEREAS**, Both parties wish to set forth the terms and conditions of their engagement;

NOW, THEREFORE, the parties agree as follows:

1. SCOPE OF SERVICES

- NovaTech shall design and implement a multi-tenant AI document ingestion pipeline using Gemini 2.5 Flash.
- Services include document classification, field extraction, human-in-the-loop triage, and full observability.
- NovaTech shall provide hybrid routing (regex fast-track for structured PDFs, GenAI for images and unstructured formats).
- Ongoing model governance and prompt versioning for twelve (12) months post-launch.
- Integration with downstream ERP/CRM systems via configurable webhooks.
- Training and documentation for up to ten (10) designated system administrators.

2. CONTRACT VALUE AND PAYMENT TERMS

The total contract value shall be **\$120,000.00 USD**, payable as follows:

| Milestone         | Amount      | Due          |
|-------------------|-------------|--------------|
| Project Kick-off  | \$30,000.00 | Feb 1, 2026  |
| MVP Delivery      | \$45,000.00 | Apr 15, 2026 |
| Production Launch | \$30,000.00 | Jun 30, 2026 |
| Final Handover    | \$15,000.00 | Dec 31, 2026 |

3. INTELLECTUAL PROPERTY

- All deliverables — system prompts, extraction schemas, pipeline architecture — remain NovaTech property until full payment.
- Upon receipt of final payment, all work product is assigned to Global Industries LLC.
- NovaTech retains the right to use anonymised learnings for product improvement.

4. CONFIDENTIALITY

- Each party shall maintain in strict confidence all Confidential Information received from the other party.
- Confidential Information includes business strategies, technical specifications, AI model configurations, and customer data.
- This obligation survives termination for three (3) years.

## **5. TECHNICAL ARCHITECTURE**

### **5.1 Document Ingestion**

- The pipeline shall accept documents via email (IMAP polling), HTTP webhook, and direct upload.
- Supported formats: PDF (native text-layer and scanned), JPEG, PNG, plain text.
- All documents are processed within fifteen (15) seconds for standard documents and sixty (60) seconds for documents exceeding ten (10) pages.

### **5.2 AI Classification**

- Document classification shall be performed using Gemini 2.5 Flash with structured JSON schema output.
- Classification shall incorporate document type descriptions, classification keywords, and field schema context.
- Classification confidence shall be reported per document and logged to the observability layer.
- The system shall support a minimum of twenty (20) document types simultaneously.

### **5.3 Field Extraction**

- Extraction shall use a single Gemini API call with a union schema covering both classification and extraction.
- For documents exceeding ten (10) pages, the system shall employ Retrieval-Augmented Generation (RAG).
- RAG implementation: chunk document into 500-token windows with 50-token overlap; embed chunks using text-embedding-004; retrieve top-K chunks per field using cosine similarity; extract from focused context.
- Extracted fields shall conform to JSON schema with type enforcement and null-safe handling.

### **5.4 Human-in-the-Loop Review**

- Documents below the configured confidence threshold shall be routed to the triage queue.
- Reviewers shall be presented with the original document, AI-extracted values, and the reason for escalation.
- Reviewer corrections shall be logged and optionally fed back as few-shot examples.
- SLA for triage review: four (4) business hours for high-priority documents.

## **6. DATA GOVERNANCE**

### **6.1 Data Residency**

- All customer data shall be processed and stored within Canada unless explicitly agreed otherwise.
- No customer document data shall be used for model training without explicit written consent.
- Data retention: run logs retained for ninety (90) days; approved records retained for seven (7) years.

### **6.2 Security**

- All API communications shall use TLS 1.3 or higher.
- Database encryption at rest using AES-256.
- API keys stored as environment variables; never logged or transmitted in plaintext.
- Access control: role-based permissions with audit logging for all administrative actions.

### **6.3 GDPR and PIPEDA Compliance**

- The pipeline shall support PII detection and redaction prior to long-term storage.
- Data subject access requests shall be fulfilled within thirty (30) calendar days.
- Privacy impact assessment shall be conducted prior to production launch.

## **7. SERVICE LEVELS**

### **7.1 Uptime**

- NovaTech shall maintain ninety-nine point five percent (99.5%) uptime measured monthly.
- Planned maintenance windows shall not exceed four (4) hours per month.
- Emergency maintenance shall be communicated with minimum two (2) hours notice.

### **7.2 Performance**

- Standard document processing: P95 latency under fifteen (15) seconds.
- Long document processing (RAG path): P95 latency under sixty (60) seconds.
- Straight-through processing rate (STP): minimum seventy-five percent (75%) within ninety (90) days of production launch.
- Triage queue clearance: all pending documents reviewed within four (4) business hours.

### **7.3 Remedies**

- For each calendar month where uptime falls below 99.5%, Client shall receive a service credit of five percent (5%) of monthly fees.
- Credits shall be applied to the following invoice and are the sole remedy for uptime failures.

## **8. AI MODEL GOVERNANCE**

### **8.1 Model Changes**

- NovaTech shall not change the AI model (e.g. from Gemini 2.5 Flash to another model) without fourteen (14) days written notice.
- All model changes shall be preceded by an evaluation run against the golden dataset.
- Evaluation results shall be shared with Client prior to deployment.

### **8.2 Prompt Versioning**

- All system prompts shall be versioned using a semantic versioning scheme (MAJOR.MINOR.PATCH).
- Prompt changes affecting extraction accuracy by more than two (2) percentage points constitute a MAJOR version change.
- Client shall be notified of all MAJOR prompt version changes with seven (7) days notice.

### **8.3 Golden Dataset**

- NovaTech shall maintain a labelled golden dataset of minimum fifty (50) documents per document type.
- The golden dataset shall be used to evaluate all configuration changes before production deployment.
- Client may contribute documents to the golden dataset subject to data governance requirements in Section 6.

### **8.4 Experiment Registry**

- All production AI configuration changes shall be recorded in the experiment registry.
- Registry entries shall include: model, prompt hash, chunking strategy, RAG configuration, confidence thresholds, and eval results.
- The experiment registry shall be available for Client inspection upon request.

## **9. SUPPORT AND MAINTENANCE**

### **9.1 Support Tiers**

- Priority 1 (System Down): Response within one (1) hour, resolution target four (4) hours.
- Priority 2 (Major Feature Impaired): Response within four (4) hours, resolution target one (1) business day.
- Priority 3 (Minor Issue): Response within one (1) business day, resolution target five (5) business days.

### **9.2 Maintenance**

- NovaTech shall provide monthly model evaluations and threshold calibration reviews.
- Quarterly business reviews to assess STP rate, cost per document, and accuracy trends.
- Annual security audit with report provided to Client within thirty (30) days of completion.

## **10. WARRANTIES AND REPRESENTATIONS**

### **10.1 NovaTech Warranties**

- NovaTech represents that it has full authority to enter into this Agreement.
- The services will be performed in a professional manner consistent with industry standards.
- The pipeline will achieve a straight-through processing rate of no less than seventy-five percent (75%) on invoice documents within ninety (90) days of production launch.
- All personnel assigned to this engagement have appropriate experience and qualifications.

### **10.2 Client Warranties**

- Client represents that it has full authority to enter into this Agreement.
- All documents provided for processing are owned by or licensed to Client.
- Client shall not use the pipeline to process documents containing illegal content.



## **11. LIMITATION OF LIABILITY**

### **11.1 Exclusion of Consequential Damages**

- In no event shall either party be liable for indirect, incidental, special, or consequential damages.
- This exclusion applies even if the party has been advised of the possibility of such damages.

### **11.2 Cap on Liability**

- NovaTech's total aggregate liability shall not exceed the total fees paid in the three (3) months preceding the claim.
- This cap applies regardless of the form of action, whether in contract, tort, or otherwise.

## **12. INDEMNIFICATION**

### **12.1 NovaTech Indemnification**

- NovaTech shall indemnify Client against claims arising from NovaTech's gross negligence or wilful misconduct.
- NovaTech shall indemnify Client against third-party IP infringement claims related to the pipeline software.

### **12.2 Client Indemnification**

- Client shall indemnify NovaTech against claims arising from Client's use of the pipeline in violation of this Agreement.
- Client shall indemnify NovaTech against claims arising from Client's document content.

## **13. DISPUTE RESOLUTION**

### **13.1 Escalation**

- Disputes shall first be escalated to senior management of both parties for good-faith negotiation.
- If not resolved within thirty (30) days, disputes shall proceed to mediation.

### **13.2 Mediation**

- Mediation shall be conducted by a mutually agreed mediator in Toronto, Ontario.
- Costs of mediation shall be shared equally between the parties.
- If mediation fails, disputes shall be resolved by binding arbitration under ICDR rules.

## **14. FORCE MAJEURE**

### **14.1 Excused Performance**

- Neither party shall be liable for delays caused by circumstances beyond their reasonable control.
- Such circumstances include: acts of God, pandemic, government action, cyberattack by state actors, or infrastructure outages.
- The affected party shall notify the other within forty-eight (48) hours of the force majeure event.

## **15. AMENDMENT AND WAIVER**

### **15.1 Amendments**

- This Agreement may only be amended by a written instrument signed by authorised representatives of both parties.
- No amendment shall be effective unless it explicitly states that it amends this Agreement.

### **15.2 Waiver**

- Failure to exercise any right under this Agreement shall not constitute a waiver of that right.
- Waivers must be in writing to be effective.

## 16. TERM AND TERMINATION

This Agreement shall commence on the Effective Date and continue for twelve (12) months. Either party may terminate with thirty (30) days written notice. In the event of material breach not cured within fifteen (15) business days, the non-breaching party may terminate immediately.

## 17. GOVERNING LAW

This Agreement shall be governed by and construed in accordance with the laws of the **Province of Ontario, Canada**, without regard to its conflict of law provisions. Any disputes shall be subject to the exclusive jurisdiction of the courts of Ontario.

## 18. ENTIRE AGREEMENT

This Agreement constitutes the entire agreement between the parties and supersedes all prior negotiations, representations, warranties, and understandings.

IN WITNESS WHEREOF, the parties have executed this Agreement as of the date first written above.

**NovaTech Solutions Ltd.**

**Global Industries LLC**

\_\_\_\_\_  
Authorized Signatory

\_\_\_\_\_  
Authorized Signatory

Date: \_\_\_\_\_

Date: \_\_\_\_\_

## SCHEDULE A — STATEMENT OF WORK

### A.1 Project Phases

- Phase 1 (Weeks 1–4): Requirements gathering, architecture design, environment setup, and proof-of-concept validation.
- Phase 2 (Weeks 5–10): Core pipeline development — ingestion, classification, extraction, and triage queue.
- Phase 3 (Weeks 11–16): Integration, testing, observability layer, and user acceptance testing.
- Phase 4 (Weeks 17–20): Production deployment, go-live support, and knowledge transfer.
- Phase 5 (Ongoing): Model governance, threshold calibration, and continuous improvement.

### A.2 Deliverables

- D1: Technical architecture document (Week 2).
- D2: Working PoC with three document types (Week 4).
- D3: MVP pipeline with full triage queue and observability (Week 10).
- D4: Production-ready system with all integrations (Week 16).

- D5: Administrator training and documentation package (Week 18).
- D6: Post-launch evaluation report including STP rate, accuracy metrics, and cost analysis (Week 24).

### **A.3 Assumptions and Dependencies**

- Client will provide sample documents for each document type within five (5) business days of project kickoff.
- Client will designate a technical point of contact available for weekly status calls.
- Client will provide access to downstream ERP/CRM systems for integration testing by Week 8.
- Gemini API access will be provisioned by NovaTech; associated API costs are included in the contract value.
- Client will complete user acceptance testing within ten (10) business days of each deliverable.
- Infrastructure will be hosted on Google Cloud Platform Canada (Montreal region).

## **SCHEDULE B — ACCEPTABLE USE POLICY**

### **B.1 Permitted Uses**

- Processing of business documents owned by or licensed to the Client.
- Integration with Client's internal systems for workflow automation.
- Training of Client staff on pipeline operation and document type configuration.

### **B.2 Prohibited Uses**

- Processing of documents containing classified government information without appropriate clearances.
- Use of the pipeline to process documents on behalf of third parties without NovaTech's written consent.
- Attempting to reverse-engineer, decompile, or extract the system prompts or AI model configuration.
- Processing documents in violation of applicable privacy laws, including PIPEDA and GDPR.
- Using the pipeline output for automated decision-making affecting individuals without human oversight.

### **B.3 Compliance Monitoring**

- NovaTech reserves the right to audit pipeline usage logs for compliance with this Acceptable Use Policy.
- Client shall cooperate fully with any compliance audit upon fourteen (14) days written notice.
- Violations of this policy may result in immediate suspension of service pending investigation.



## **SCHEDULE C — DATA PROCESSING AGREEMENT**

### **C.1 Roles and Responsibilities**

- For the purposes of PIPEDA and applicable privacy legislation, Client is the data controller and NovaTech is the data processor.
- NovaTech shall process personal data only on documented instructions from Client.
- NovaTech shall ensure that persons authorised to process personal data are bound by confidentiality obligations.

### **C.2 Sub-processors**

- NovaTech uses Google LLC as a sub-processor for AI inference (Gemini API).
- Client consents to NovaTech's use of Google LLC as a sub-processor for the purposes of this Agreement.
- NovaTech shall notify Client of any intended changes to sub-processors with thirty (30) days notice.
- Any new sub-processor must meet equivalent data protection standards to those in this Schedule.

### **C.3 Data Subject Rights**

- NovaTech shall assist Client in fulfilling data subject access requests within five (5) business days of receipt.
- NovaTech shall notify Client within forty-eight (48) hours of becoming aware of a personal data breach.
- Upon termination, NovaTech shall delete all personal data within thirty (30) days unless legally required to retain it.

### **C.4 International Data Transfers**

- NovaTech shall not transfer personal data outside Canada without Client's prior written consent.
- Where transfers are necessary (e.g. Gemini API processing), NovaTech shall ensure adequate safeguards are in place.
- Standard contractual clauses or equivalent mechanisms shall be used for any cross-border transfers.

## **SCHEDULE D — CHANGE MANAGEMENT PROCESS**

### **D.1 Change Request Procedure**

- Either party may submit a Change Request (CR) using the agreed CR template.
- NovaTech shall assess the impact of the CR within five (5) business days and provide a written estimate.
- Changes shall only proceed upon written acceptance of the estimate by Client.
- Emergency changes required to maintain service availability may be implemented immediately with retroactive approval.

### **D.2 Change Categories**

- Category 1 (Minor): Changes with no impact on cost, timeline, or service levels. No CR required.
- Category 2 (Standard): Changes with minor cost or timeline impact. CR required, approved by project managers.
- Category 3 (Major): Changes with significant cost, timeline, or scope impact. CR required, approved by executives.

### **D.3 AI Configuration Changes**

- Changes to document type definitions, field schemas, or confidence thresholds by Client are Category 1.
- Changes to system prompts, model selection, or RAG configuration by NovaTech are Category 2.
- Introduction of new AI capabilities or architectural changes are Category 3.

## **SCHEDULE E — PRICING AND INVOICING**

### **E.1 Fee Structure**

- The contract value of \$120,000.00 USD covers all services described in Schedule A for a period of twelve (12) months.
- AI inference costs (Gemini API) are included in the contract value up to a monthly cap of \$500.00 USD.
- Inference costs exceeding \$500.00 USD per month shall be invoiced separately at cost plus ten percent (10%).
- Additional document types beyond the initial four (4) shall be billed at \$2,000.00 USD per type.

### **E.2 Invoicing**

- Invoices shall be issued on the dates specified in the payment schedule in Section 2.
- Payment terms: Net 30 from invoice date.
- Late payments shall accrue interest at one and a half percent (1.5%) per month.
- NovaTech may suspend service for invoices unpaid after sixty (60) days.

### **E.3 Expense Reimbursement**

- Pre-approved travel expenses shall be reimbursed at cost.
- All reimbursable expenses must be submitted with receipts within thirty (30) days of incurrence.
- Expenses exceeding \$1,000.00 CAD require prior written approval from Client.

## SCHEDULE F — TECHNICAL SPECIFICATIONS

### F.1 System Architecture

- Frontend: Streamlit web application hosted on Google Cloud Run.
- Backend: Python 3.12 with FastAPI for webhook endpoints.
- Database: SQLite for PoC; PostgreSQL for production multi-tenant deployment.
- AI: Gemini 2.5 Flash via Google Generative Language API (v1beta).
- Embeddings: text-embedding-004 for RAG vector index.
- Vector store: FAISS in-memory for PoC; Vertex AI Vector Search for production.

### F.2 API Specifications

- Classification endpoint: POST /api/v1/classify — accepts multipart/form-data with document and doc\_type\_hint.
- Extraction endpoint: POST /api/v1/extract — accepts document and confirmed doc\_type.
- Webhook endpoint: POST /api/v1/webhook — accepts JSON payload with document URL and metadata.
- All endpoints return JSON with fields: status, doc\_type, confidence, extracted\_fields, route, latency\_ms, cost\_usd.
- Rate limits: 100 requests per minute per tenant; 1000 requests per hour.

### F.3 Integration Requirements

- ERP integration via REST webhook — Client provides target URL and authentication credentials.
- Email ingestion via IMAP polling — Client provides mailbox credentials and folder structure.
- SSO integration via SAML 2.0 or OAuth 2.0 — Client provides IdP metadata.
- Audit log export: SIEM-compatible JSON format, exportable via API or scheduled delivery.

## **SCHEDULE G — DISASTER RECOVERY AND BUSINESS CONTINUITY**

### **G.1 Recovery Objectives**

- Recovery Time Objective (RTO): Pipeline restored within four (4) hours of a major outage.
- Recovery Point Objective (RPO): Maximum data loss of one (1) hour.
- All approved records are replicated to a secondary region within fifteen (15) minutes.

### **G.2 Backup Procedures**

- Database backups: every six (6) hours, retained for thirty (30) days.
- Configuration backups: on every change, retained indefinitely.
- Document store backups: daily, retained for seven (7) years in compliance with financial record requirements.
- Backup restoration is tested quarterly; test results are available to Client upon request.

### **G.3 Incident Response**

- NovaTech maintains a documented incident response plan reviewed annually.
- Security incidents are classified as P1 (Critical), P2 (High), or P3 (Medium).
- P1 incidents trigger immediate escalation to the Client's designated security contact.
- Post-incident reports are delivered within five (5) business days of resolution.

## **SCHEDULE H — ACCEPTANCE TESTING CRITERIA**

### **H.1 Performance Acceptance Criteria**

- Straight-through processing rate  $\geq 75\%$  on a representative set of 100 invoice documents.
- P95 classification latency  $< 5$  seconds for image documents.
- P95 extraction latency  $< 10$  seconds for standard documents ( $< 10$  pages).
- P95 extraction latency  $< 60$  seconds for long documents ( $\geq 10$  pages, RAG path).
- Zero data loss: all documents processed are either approved, triaged, or marked as failed — none silently dropped.

### **H.2 Accuracy Acceptance Criteria**

- Field-level precision  $\geq 90\%$  on the golden dataset for Invoice and Purchase Order types.
- Field-level precision  $\geq 85\%$  on the golden dataset for Complaint and Contract types.
- Confidence calibration: documents scored  $\geq 0.90$  confidence shall have  $\geq 95\%$  actual field accuracy.
- False positive rate (wrong data auto-approved):  $< 1\%$  on the golden dataset.

### **H.3 Functional Acceptance Criteria**

- All four document types correctly classified from a mixed set of 50 documents.
- Triage queue correctly populated for all documents below the configured threshold.
- Reviewer corrections successfully saved and reflected in approved records.
- Observability log captures route, cost, latency, and AI prompts/responses for 100% of runs.
- Webhook integration successfully pushes approved records to a test ERP endpoint.

### **H.4 Acceptance Process**

- Client shall conduct UAT within ten (10) business days of each deliverable.
- Defects are classified as Critical (blocks acceptance), Major (must fix before go-live), or Minor (fix in next release).
- Critical and Major defects must be resolved before acceptance is granted.
- Acceptance is granted in writing by Client's designated technical lead.

## **SCHEDULE I — KNOWLEDGE TRANSFER**

### **I.1 Documentation**

- Architecture overview: system components, data flows, and integration points.
- Operator guide: model governance procedures, prompt versioning, and eval runner usage.
- Administrator guide: document type configuration, field definitions, and sensitivity settings.
- User guide: document upload, review queue operation, and observability dashboard.
- API reference: all endpoints, request/response schemas, and error codes.
- Runbook: common operational procedures including backup restoration and incident response.

### **I.2 Training**

- Session 1 (4 hours): System overview and architecture walkthrough for technical team.
- Session 2 (4 hours): Administrator training — document type configuration and AI governance.
- Session 3 (2 hours): End-user training — document processing and review queue.
- Session 4 (2 hours): Operator training — model governance, eval runner, and incident response.
- All sessions recorded and provided to Client for future onboarding.

### **I.3 Hypercare Period**

- NovaTech shall provide dedicated support for thirty (30) days post go-live.
- Hypercare includes daily status calls, priority response to all issues, and proactive monitoring.
- Any defects discovered during hypercare are remediated at no additional cost.

## **SCHEDULE J — ENVIRONMENTAL AND SOCIAL RESPONSIBILITY**

### **J.1 Carbon Footprint**

- NovaTech commits to measuring the carbon footprint of AI inference for this engagement.
- Google Cloud Platform's carbon-neutral operations are used where available.
- Annual sustainability report provided to Client including estimated CO2e for AI processing.

### **J.2 Ethical AI**

- NovaTech commits to responsible AI practices including transparency, fairness, and human oversight.
- The human-in-the-loop triage queue ensures no fully automated decision-making for high-stakes documents.
- Bias monitoring: NovaTech will evaluate extraction accuracy across document variants to identify systematic biases.
- Client retains full control over confidence thresholds and review sensitivity — no AI-only approvals without explicit opt-in.



## **SCHEDULE K — NOTICES**

### **K.1 Notice Requirements**

- All notices under this Agreement shall be in writing and delivered by email with read receipt or by courier.
- Notices to NovaTech: legal@novatech-solutions.com with copy to ceo@novatech-solutions.com.
- Notices to Client: legal@globalindustries.com with copy to operations@globalindustries.com.
- Notices are effective on the date of confirmed receipt.

## SCHEDULE L — DEFINITIONS

### L.1 Key Definitions

- "Approved Record" means a document whose extracted fields have been confirmed by either auto-approval or human review.
- "Classification" means the process of identifying the document type of an ingested document.
- "Confidence Score" means the AI-reported probability (0.0 to 1.0) that extracted fields are accurate.
- "Document Type" means a configured category of document with defined fields and classification parameters.
- "Extraction" means the process of identifying and returning specified field values from a classified document.
- "Fast-Track" means deterministic field extraction using anchor keywords and regex patterns on PDF text layers.
- "Golden Dataset" means a curated set of labelled documents used to evaluate AI configuration changes.
- "Human-in-the-Loop" means a workflow step requiring human review and approval before a record is finalised.
- "RAG" means Retrieval-Augmented Generation — a technique that retrieves relevant document sections before AI extraction.
- "Straight-Through Processing (STP)" means automated document processing without human review.
- "Triage Queue" means the queue of documents requiring human review before approval.

## **SCHEDULE M — AUDIT RIGHTS**

### **M.1 NovaTech Audit Rights**

- NovaTech reserves the right to audit Client's use of the pipeline to ensure compliance with this Agreement and the Acceptable Use Policy.
- Audits shall be conducted no more than once per year unless a compliance concern is identified.
- Client shall provide reasonable cooperation and access to relevant records.
- Audit findings shall be shared with Client within ten (10) business days of completion.

### **M.2 Client Audit Rights**

- Client may audit NovaTech's security controls, data handling procedures, and AI governance processes once per year.
- NovaTech shall provide SOC 2 Type II report or equivalent upon request.
- Third-party audits requested by Client shall be at Client's cost.
- NovaTech shall remediate any audit findings categorised as High or Critical within sixty (60) days.

## **SCHEDULE N — INSURANCE**

### **N.1 Required Coverage**

- NovaTech shall maintain the following insurance throughout the term:
- Commercial General Liability: minimum \$2,000,000 per occurrence.
- Professional Liability (Errors and Omissions): minimum \$2,000,000 per claim.
- Cyber Liability: minimum \$5,000,000 per occurrence.
- Workers Compensation: as required by applicable law.
- Evidence of insurance shall be provided to Client within five (5) business days of request.

## **SCHEDULE O — SUBCONTRACTING**

### **O.1 Approved Subcontractors**

- NovaTech may subcontract portions of the services to pre-approved subcontractors.
- Approved subcontractors as of the Effective Date: Google LLC (AI inference), Amazon Web Services (backup storage).
- NovaTech shall ensure all subcontractors are bound by confidentiality and data protection obligations equivalent to this Agreement.
- NovaTech remains fully responsible for the performance of all subcontractors.

### **O.2 New Subcontractors**

- NovaTech shall provide thirty (30) days written notice before engaging any new subcontractor.
- Client may object to a new subcontractor on reasonable grounds within fifteen (15) days of notice.
- If Client objects, NovaTech shall either not engage the subcontractor or propose an acceptable alternative.

## **SCHEDULE P — TRANSITION ASSISTANCE**

### **P.1 End-of-Term Transition**

- Upon expiry or termination, NovaTech shall provide up to sixty (60) days of transition assistance.
- Transition assistance includes: data export in standard formats, documentation of all configurations, and knowledge transfer to Client or successor provider.
- Transition assistance is billed at NovaTech's then-current professional services rate.

### **P.2 Data Return**

- Within thirty (30) days of termination, NovaTech shall provide Client with a complete export of all approved records, run logs, and configuration data.
- Data shall be delivered in JSON and CSV formats.
- NovaTech shall delete all Client data within sixty (60) days of the export, subject to legal retention requirements.

### **P.3 Continuity**

- NovaTech shall cooperate in good faith with any successor provider to ensure service continuity.
- API documentation and integration specifications shall be provided to the successor provider upon Client's written instruction.

## EXHIBIT 1 — DOCUMENT TYPE REFERENCE

### Invoice

Billing document from vendor to client. Fields: invoice\_number, customer\_name, total\_amount.

Keywords: INVOICE, TAX INVOICE. Sensitivity: Medium (85% threshold).

### Purchase Order

Authorization to purchase goods/services. Fields: po\_number, ship\_to\_company, quantity. Keywords:

PURCHASE ORDER, REQUISITION. Sensitivity: Medium (85% threshold).

### Complaint

Customer support escalation. Fields: ticket\_id, customer\_name, issue\_summary. Keywords:

COMPLAINT, GRIEVANCE. Sensitivity: High (92% threshold).

### Contract

Legal agreement between parties. Fields: contract\_id, parties, effective\_date, governing\_law,

contract\_value. Keywords: AGREEMENT, WHEREAS. Sensitivity: High (92% threshold).

## EXHIBIT 2 — CONFIDENCE CALIBRATION REFERENCE

The following table reflects observed correlation between Gemini confidence scores and actual field extraction accuracy, measured on the golden dataset. Thresholds are set conservatively to prioritise precision over recall — a false positive (wrong data approved) is more costly than a false negative (correct data sent to triage).

| Confidence Band | Invoice Accuracy | Complaint Accuracy | Contract Accuracy |
|-----------------|------------------|--------------------|-------------------|
| 0.50 – 0.60     | 72%              | 45%                | 51%               |
| 0.60 – 0.70     | 81%              | 58%                | 63%               |
| 0.70 – 0.80     | 88%              | 67%                | 71%               |
| 0.80 – 0.90     | 94%              | 74%                | 79%               |
| 0.90 – 1.00     | 98%              | 87%                | 91%               |

Note: These figures are derived from testing on the initial golden dataset. Accuracy improves over time as reviewer corrections are fed back as few-shot examples and thresholds are recalibrated. The calibration table in Exhibit 2 shall be updated quarterly as part of the ongoing governance process.